

## Bells and whistles in neural net training



## Tricks in training neural networks

There are various tricks that people use when training neural networks:

- ▶ Regularization: Adjusting the gradient
- ▶ Dropout: Adjusting the hidden units
- ▶ Optimization methods: Adjusting the learning rate
- ▶ Initialization: Using particular forms of initialization

## Regularization

Neural networks can be regularized in a similar way as linear models. Neural networks can also with **Frobenius norm**, which is a trivial extension to L2 norm for matrices. In fact, in many cases it is just referred to as L2 regularization.

$$\mathcal{L} = \sum_{i=1}^N \ell^{(i)} + \lambda_{z \rightarrow y} \|\Theta^{(z \rightarrow y)}\|_F^2 + \lambda_{x \rightarrow z} \|\Theta^{(x \rightarrow z)}\|_F^2$$

where  $\|\Theta\|_F^2 = \sum_{i,j} \theta_{i,j}^2$  is the squared **Frobenius norm**, which generalizes the  $L_2$  norm to matrices. The bias parameters  $b$  are not regularized, as they do not contribute to the classifier to the inputs.

## L2 regularization

- Compute the gradient of a loss with L2 regularization

$$\frac{\partial \mathcal{L}}{\partial \theta} = \sum_{i=1}^N \frac{\partial \ell^{(i)}}{\partial \theta} + \lambda \theta$$

Assignment Project Exam Help

- Update the weights <https://powcoder.com>

Add WeChat powcoder

$$\theta = \theta - \eta \left( \sum_{i=1}^N \frac{\partial \ell^{(i)}}{\partial \theta} + \lambda \theta \right)$$

- “Weigh decay factor”:  $\lambda$  is a tunable hyper parameter that pulls a weight back when it has become too big
- Question: Does it matter which layer  $\theta$  is from when computing the regularization term?

## L1 regularization

- L1 regularization loss

$$\mathcal{L} = \sum_{i=1}^N \ell^{(i)} + \lambda_{z \rightarrow y} \|\Theta^{(z \rightarrow y)}\|_1 + \lambda_{x \rightarrow z} \|\Theta^{(x \rightarrow z)}\|_1$$

Assignment Project Exam Help

- Compute the gradient

<https://powcoder.com>

$$\frac{\partial \mathcal{L}}{\partial \theta} = \sum_{i=1}^N \frac{\partial \ell^{(i)}}{\partial \theta} + \lambda \operatorname{sign}(\theta)$$

Add WeChat powcoder

- update the weights

$$\theta = \theta - \eta \left( \sum_{i=1}^N \frac{\partial \ell^{(i)}}{\partial \theta} + \lambda \operatorname{sign}(\theta) \right)$$

## Comparison of L1 and L2

- ▶ In L1 regularization, the weights shrink by a constant amount toward 0. In L2 regularization, the weights shrink by an amount which is proportional to  $w$ .
- ▶ When a particular weight has a large absolute value,  $|\theta|$ , L1 regularization shrinks the weight much less than L2 regularization does. By contrast, when  $|\theta|$  is small, L1 regularization shrinks the weight much more than L2 regularization.
- ▶ The net result is that L1 regularization tends to concentrate the weight of the network in a relatively small number of high-importance connections, while the other weights are driven toward zero. So L1 regularization effectively does *feature selection*.

## Dropout

- ▶ Randomly drops a certain percentage of the nodes to prevent over-reliance on a few features or hidden units, or **feature co-adaptation**, where some features are only useful when working together with a few other features. The ultimate goal is to avoid overfitting.

## Dropout

- Dropout can be achieved using a mask:

$$\mathbf{z}^{(1)} = g^1(\Theta^{(1)}\mathbf{x} + \mathbf{b}^1)$$

$$\mathbf{m}^1 \sim \text{Bernoulli}(r^1)$$

$$\tilde{\mathbf{z}}^{(1)} = \mathbf{m}^1 \odot \mathbf{z}^{(1)}$$

$$\mathbf{z}^{(2)} = g^2(\Theta^{(2)}\tilde{\mathbf{z}}^{(1)} + \mathbf{b}^2)$$

$$\mathbf{m}^2 \sim \text{Bernoulli}(r^2)$$

$$\tilde{\mathbf{z}}^{(2)} = \mathbf{m}^2 \odot \mathbf{z}^{(2)}$$

$$\mathbf{y} = \Theta^{(3)}\tilde{\mathbf{z}}^{(2)}$$

where  $\mathbf{m}^1$  and  $\mathbf{m}^2$  are mask vectors. The values of the elements in these vectors are either 1 or 0, drawn from a Bernoulli distribution with parameter  $r$  (usually  $r = 0.5$ )



## Optimization methods

- ▶ SGD with Momentum
  - ▶ AdaGrad
  - ▶ Root Mean Square Prop (RMSProp)
  - ▶ Adam
- Assignment Project Exam Help  
<https://powcoder.com>  
Add WeChat powcoder

## SGD with Momentum

- At each timestep  $t$ , compute  $\nabla_{\theta}\mathcal{L}$ , and then compute the momentum as follows:

Assignment Project Exam Help

$$V_0 = 0, \beta \approx 0.9$$
$$V_t = \beta V_{t-1} + (1 - \beta) \nabla_{\theta} \mathcal{L}$$

Add WeChat powcoder

- The momentum term increases for dimensions whose gradient point in the same directions and reduces updates for dimensions whose gradient change directions.

## AdaGrad

- Keep a running sum of the squared gradient  $V_{\nabla_{\theta}}$ . When updating the weight of this *theta*, divide the gradient by the square root of this term

Assignment Project Exam Help

$V_t = V_{t-1} + \nabla_{\theta} \mathcal{L}^2$   
<https://powcoder.com>

$\theta_j = \theta_j - \eta \frac{\nabla_{\theta} \mathcal{L}}{\sqrt{V_t} + \epsilon}$   
Add WeChat powcoder

e.g.,  $\epsilon = 10^{-8}$

- The net effect is to slow down the update for weights with large gradient and accelerate the update for weights with small gradient

## Root Mean Square Prop (RMSProp)

- A minor adjustment of AdaGrad. Instead of letting the sum of squared gradient continuously grow, we let the sum decay:

Assignment Project Exam Help

$$V_0 = 0$$

$$V_t = \beta V_{t-1} + (1 - \beta) \nabla_{\theta} \mathcal{L}^2$$

$$\theta_j = \theta_j - \eta \frac{\nabla_{\theta} \mathcal{L}}{\sqrt{V_t} + \epsilon}$$

Add WeChat powcoder

e.g.  $\beta \approx 0.9, \eta = 0.001, \epsilon = 10^{-8}$

## Adaptive Moment Estimation (Adam)

- Weight update at time step  $t$  for Adam:

$$V_0 = 0, S_0 = 0,$$

$$V_t = \beta_1 V_{t-1} + (1 - \beta_1) \nabla_{\theta} \mathcal{L} \quad \text{Momentum}$$

$$S_t = \beta_2 S_{t-1} + (1 - \beta_2) \nabla_{\theta} \mathcal{L}^2 \quad \text{RMSProp}$$

$$V_t^{\text{corrected}} = \frac{V_t}{\beta_1^t}$$

$$S_t^{\text{corrected}} = \frac{S_t}{\beta_2^t}$$

$$\theta_j = \theta_j - \eta \frac{V_t^{\text{corrected}}}{\sqrt{S_t^{\text{corrected}} + \epsilon}}$$

- Adam combines Momentum and RMSProp

## Initialization

Xavier Initialization:

Assignment Project Exam Help

$$\Theta \sim \mathcal{U}\left[-\sqrt{\frac{6}{n^{(l)} + n^{(l+1)}}}, \sqrt{\frac{6}{n^{(l)} + n^{(l+1)}}}\right]$$

where  $n^{(l)}$  is the number of input units to  $\Theta$  (fan-in),  $n^{(l+1)}$  is the number of output units from  $\Theta$

Add WeChat powcoder

## Neural net in PyTorch

```
from torch import nn
class Net(nn.Module):
    """subclass from nn.Module is
    important to inspecting the parameters"""

    def __init__(self, in_dim=25, out_dim=3, batch_size=1):
        super(Net, self).__init__()
        self.in_dim = in_dim
        self.out_dim = out_dim
        self.linear = nn.Linear(self.in_dim, self.out_dim)
        self.softmax = nn.Softmax(dim=1)

    def forward(self, input_matrix):
        logit = self.linear(input_matrix)
        #return raw score, not normalized score
        return logit

    def xentropy_loss(self, input_matrix, target_label_vec):
        loss = nn.CrossEntropyLoss()
        logits = self.forward(input_matrix)
        return loss(logits, target_label_vec)
```

## Use optimizers in Pytorch

```
import torch.optim as optim
net = Net(input_dim, output_dim)
optimizer = optim.Adam(net.parameters(), lr=lr)
for epoch in range(epochs):
    total_nll = 0
    for batch in batchify(train_data, batch_size):
        optimizer.zero_grad() #zero out the gradient.
        vectorized = vectorize_batch(batch, \
                                     feat_index, label_index)
        feat_vec = map(itemgetter(0), vectorized)
        label_vec = map(itemgetter(1), vectorized)
        feat_list = list(feat_vec)
        label_list = list(label_vec)
        x = torch.Tensor(feat_list)
        y = torch.LongTensor(label_list)
        loss = net.xentropy_loss(x, y)
        total_nll += loss
        loss.backward()
        optimizer.step()
    torch.save(net.state_dict(), net_path)
```